

Stochastic Program Optimization

Eric Schkufza, Rahul Sharma, and Alex Aiken, CACM'16

Presented by Ajay Jain

Reading group | April 4th, 2018

Program synthesis & induction

- Target program: $\vec{y} = P(\vec{x})$
- Generated program: $\hat{P}(\vec{x}) \approx \vec{y}$
- Generated using:
 - input output pairs
 - an example program P (e.g. for performance optimization or translation/compilation)
- In **program synthesis**, \hat{P} is interpretable
 - e.g. stored as explicit code
- In **program induction**, \hat{P} can be evaluated, but we don't know its internals

Program synthesis & induction **for compilers**

- Compilers need to be **correct** (mostly - according to a spec) in translating a program to output machine instructions
- Typically, you hope to **minimize runtime**
- In a reasonable amount of time (compilation is already pretty slow...)

Superoptimization

Figure 10. SAXPY benchmark.

```
void SAXPY(int* x, int* y, int a) {
    x[i]  = a * x[i]  + y[i];
    x[i+1] = a * x[i+1] + y[i+1];
    x[i+2] = a * x[i+2] + y[i+2];
    x[i+3] = a * x[i+3] + y[i+3];
}

1 # gcc -O3                1 # STOKE
2                          2
3 movslq ecx,rcx           3 movd edi,xmm0
4 leaq (rsi,rcx,4),r8      4 shufps 0,xmm0,xmm0
5 leaq 1(rcx),r9           5 movups (rsi,rcx,4),xmm1
6 movl (r8),eax           6 pmullw xmm1,xmm0
7 imull edi,eax           7 movups (rdx,rcx,4),xmm1
8 addl (rdx,rcx,4),eax     8 paddw xmm1,xmm0
9 movl eax,(r8)           9 movups xmm0,(rsi,rcx,4)
10 leaq (rsi,r9,4),r8
11 movl (r8),eax
12 imull edi,eax
13 addl (rdx,r9,4),eax
14 leaq 2(rcx),r9
15 addq 3,rcx
16 movl eax,(r8)
17 leaq (rsi,r9,4),r8
18 movl (r8),eax
19 imull edi,eax
20 addl (rdx,r9,4),eax
21 movl eax,(r8)
22 leaq (rsi,rcx,4),rax
23 imull (rax),edi
24 addl (rdx,rcx,4),edi
25 movl edi,(rax)
```

- **Basic block:** A sequence of instructions without branches (jumps).
 - Also known as straight-line code
- **Superoptimization:** Translate straight-line code into a more optimized form



Superoptimization

Figure 10. SAXPY benchmark.

```
void SAXPY(int* x, int* y, int a) {
    x[i]  = a * x[i]  + y[i];
    x[i+1] = a * x[i+1] + y[i+1];
    x[i+2] = a * x[i+2] + y[i+2];
    x[i+3] = a * x[i+3] + y[i+3];
}

1 # gcc -O3                1 # STOKE
2                          2
3 movslq ecx,rcx           3 movd edi,xmm0
4 leaq (rsi,rcx,4),r8      4 shufps 0,xmm0,xmm0
5 leaq 1(rcx),r9           5 movups (rsi,rcx,4),xmm1
6 movl (r8),eax           6 pmullw xmm1,xmm0
7 imull edi,eax           7 movups (rdx,rcx,4),xmm1
8 addl (rdx,rcx,4),eax     8 paddw xmm1,xmm0
9 movl eax,(r8)           9 movups xmm0,(rsi,rcx,4)
10 leaq (rsi,r9,4),r8
11 movl (r8),eax
12 imull edi,eax
13 addl (rdx,r9,4),eax
14 leaq 2(rcx),r9
15 addq 3,rcx
16 movl eax,(r8)
17 leaq (rsi,r9,4),r8
18 movl (r8),eax
19 imull edi,eax
20 addl (rdx,r9,4),eax
21 movl eax,(r8)
22 leaq (rsi,rcx,4),rax
23 imull (rax),edi
24 addl (rdx,rcx,4),edi
25 movl edi,(rax)
```

The guy gcc
tells you not
to worry
about

You

- **Basic block:** A sequence of instructions without branches (jumps).
 - Also known as straight-line code
- **Superoptimization:** Translate straight-line code into a more optimized form



To be fair to gcc...

Figure 10. SAXPY benchmark.

```
void SAXPY(int* x, int* y, int a) {
    x[i]  = a * x[i]  + y[i];
    x[i+1] = a * x[i+1] + y[i+1];
    x[i+2] = a * x[i+2] + y[i+2];
    x[i+3] = a * x[i+3] + y[i+3];
}

1 # gcc -O3                1 # STOKE
2                          2
3 movslq ecx,rcx           3 movd edi,ymm0
4 leaq (rsi,rcx,4),r8      4 shufps 0,ymm0,ymm0
5 leaq 1(rcx),r9           5 movups (rsi,rcx,4),ymm1
6 movl (r8),eax           6 pmullw ymm1,ymm0
7 imull edi,eax           7 movups (rdx,rcx,4),ymm1
8 addl (rdx,rcx,4),eax     8 paddw ymm1,ymm0
9 movl eax,(r8)           9 movups ymm0,(rsi,rcx,4)
10 leaq (rsi,r9,4),r8
11 movl (r8),eax
12 imull edi,eax
13 addl (rdx,r9,4),eax
14 leaq 2(rcx),r9
15 addq 3,rcx
16 movl eax,(r8)
17 leaq (rsi,r9,4),r8
18 movl (r8),eax
19 imull edi,eax
20 addl (rdx,r9,4),eax
21 movl eax,(r8)
22 leaq (rsi,rcx,4),rax
23 imull (rax),edi
24 addl (rdx,rcx,4),edi
25 movl edi,(rax)
```

```
10 void SAXPY_noalias_alt(int * restrict out, int
11 * restrict x, int * restrict y, int a) {
12     out[i+0] = a * x[i+0] + y[i+0];
13     out[i+1] = a * x[i+1] + y[i+1];
14     out[i+2] = a * x[i+2] + y[i+2];
15     out[i+3] = a * x[i+3] + y[i+3];
16 }

17 SAXPY_noalias_alt:
18 .LFB2:
19     .cfi_startproc
20     movslq i(%rip), %rax
21     movl %ecx, -12(%rsp)
22     movd -12(%rsp), %xmm3
23     pshufd $0, %xmm3, %xmm0
24     leaq 0(,%rax,4), %r8
25     movdqu (%rsi,%r8), %xmm2
26     movdqa %xmm2, %xmm1
27     psrlq $32, %xmm2
28     pmuludq %xmm0, %xmm1
29     psrlq $32, %xmm0
30     pshufd $8, %xmm1, %xmm1
31     pmuludq %xmm2, %xmm0
32     pshufd $8, %xmm0, %xmm0
33     punpckldq %xmm0, %xmm1
34     movdqa %xmm1, %xmm0
35     movdqu (%rdx,%r8), %xmm1
36     padd %xmm1, %xmm0
37     movups %xmm0, (%rdi,%rax,4)
38     ret
```

- Originally, gcc translated the SAXPY function into a series of *scalar operations* for safety, possibly due to aliasing and alignment issues. Annotations help.



Cost to minimize

$$\text{cost}(\mathcal{R}; \mathcal{T}) = w_e \cdot \underset{\text{Correctness}}{\text{eq}(\mathcal{R}; \mathcal{T})} + w_p \cdot \underset{\text{Performance}}{\text{perf}(\mathcal{R})} \quad (1)$$

- **eq(Rewrite; Target):**
measures similarity
 $\text{eq}(\mathcal{R}, \mathcal{T}) = 0$ if \mathcal{R} and \mathcal{T} result
in the same output registers
and memory given the same
input (**live inputs and outputs*)

Cost to minimize

$$\text{cost}(\mathcal{R}; \mathcal{T}) = w_e \cdot \underset{\text{Correctness}}{\text{eq}(\mathcal{R}; \mathcal{T})} + w_p \cdot \underset{\text{Performance}}{\text{perf}(\mathcal{R})} \quad (1)$$

$$\text{eq}^*(\mathcal{R}; \mathcal{T}, \tau) = \begin{cases} \text{eq}(\mathcal{R}; \mathcal{T}), & \text{if } \text{eq}'(\mathcal{R}; \mathcal{T}, \tau) = 0 \\ \text{eq}'(\mathcal{R}; \mathcal{T}, \tau), & \text{otherwise} \end{cases}$$

- **eq(Rewrite; Target):**
measures similarity
 $\text{eq}(\mathcal{R}, \mathcal{T}) = 0$ if \mathcal{R} and \mathcal{T} result
in the same output registers
and memory given the same
input (**live inputs and outputs*)
 - *With a symbolic validator:* only 1000
evaluations per second
 - *→ Solution:* approximate based on test cases

Cost to minimize

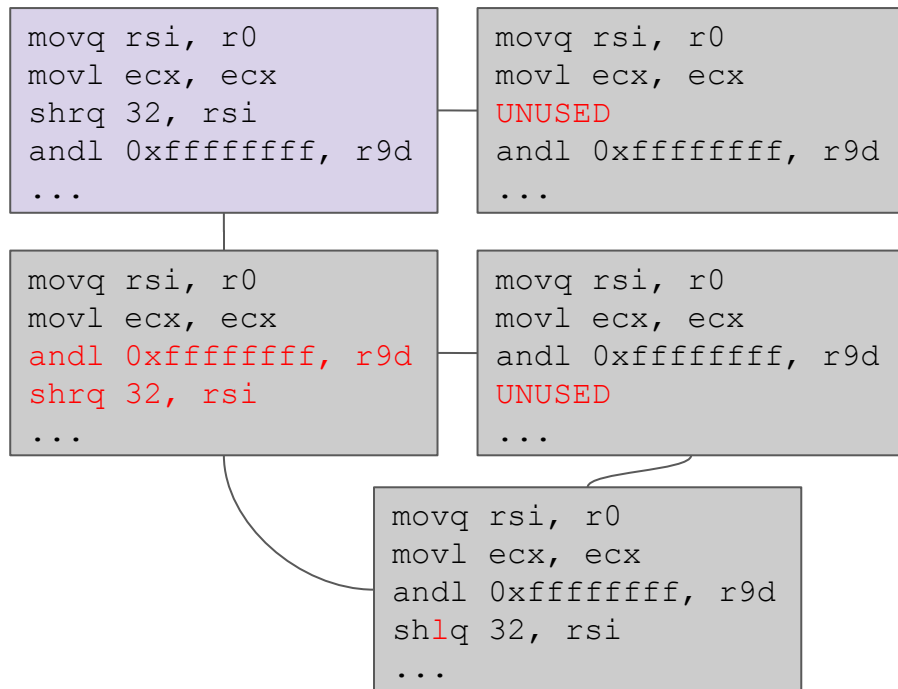
$$\text{cost}(\mathcal{R}; \mathcal{T}) = w_e \cdot \underset{\text{Correctness}}{\text{eq}(\mathcal{R}; \mathcal{T})} + w_p \cdot \underset{\text{Performance}}{\text{perf}(\mathcal{R})} \quad (1)$$

$$\text{eq}^*(\mathcal{R}; \mathcal{T}, \tau) = \begin{cases} \text{eq}(\mathcal{R}; \mathcal{T}), & \text{if } \text{eq}'(\mathcal{R}; \mathcal{T}, \tau) = 0 \\ \text{eq}'(\mathcal{R}; \mathcal{T}, \tau), & \text{otherwise} \end{cases}$$

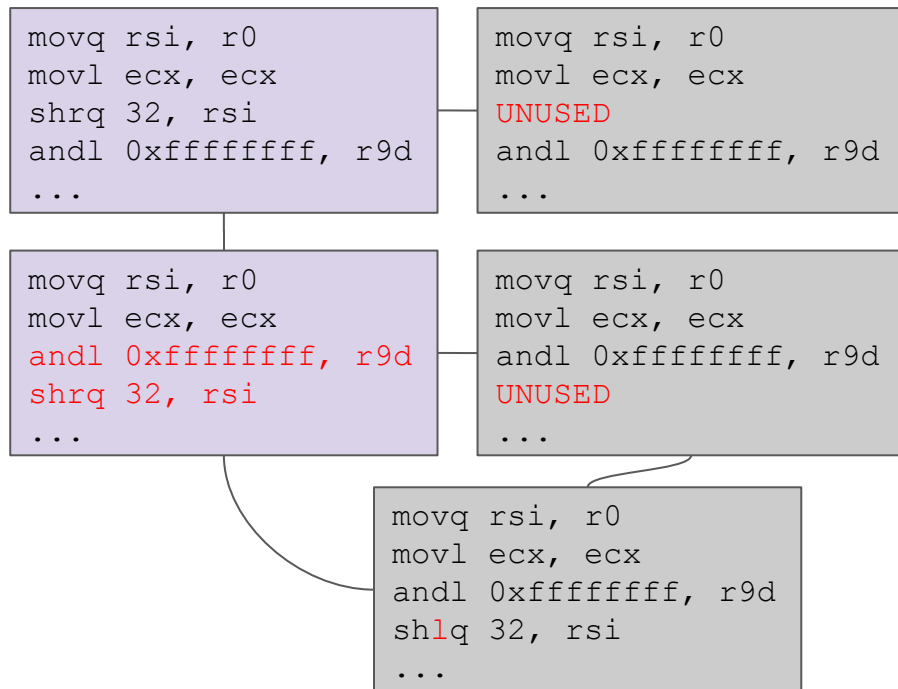
$$\text{perf}(\mathcal{R}) = \sum_{i \in \text{inst}(\mathcal{R})} \text{lat}(i)$$

- **eq(Rewrite; Target):**
measures similarity
 $\text{eq}(\mathcal{R}, \mathcal{T}) = 0$ if \mathcal{R} and \mathcal{T} result in the same output registers and memory given the same input (**live inputs and outputs*)
 - With a symbolic validator: only 1000 evaluations per second
 - → *Solution*: approximate based on test cases
- **perf:** Quantifies speedup of the rewrite
 - (heuristic, sum of instruction latencies)

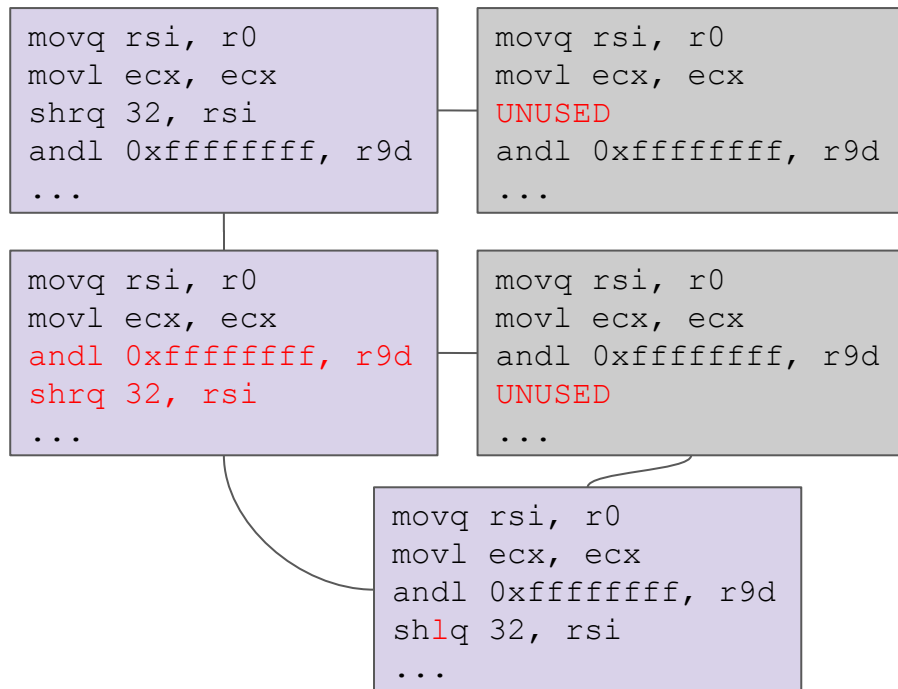
Search space - Markov Chain, graph for program space



Search space - Markov Chain, graph for program space



Search space - Markov Chain, graph for program space



Search space - Markov Chain, graph for program space

- **Opcode.** An instruction is randomly selected, and its opcode is replaced by a random opcode.
- **Operand.** An instruction is randomly selected and one of its operands is replaced by a random operand.
- **Swap.** Two lines of code are randomly selected and interchanged.
- **Instruction.** An instruction is randomly selected and replaced either by a random instruction or the UNUSED token. Proposing UNUSED corresponds to deleting an instruction, and replacing UNUSED by an instruction corresponds to inserting an instruction.

```
movq rsi, r0
movl ecx, ecx
shrq 32, rsi
andl 0xffffffff, r9d
...
```

```
movq rsi, r0
movl ecx, ecx
UNUSED
andl 0xffffffff, r9d
...
```

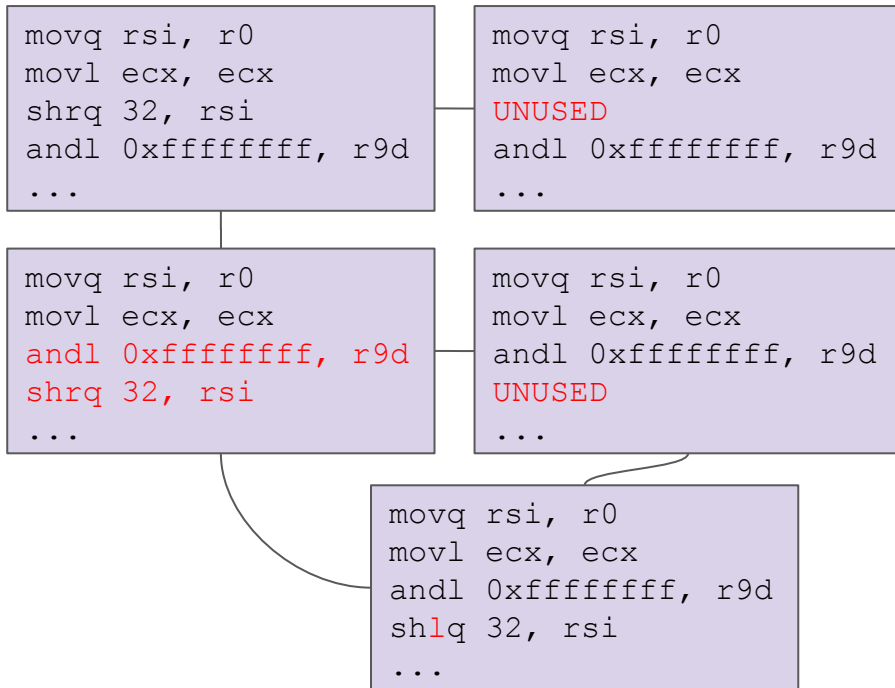

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
shrq 32, rsi
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
UNUSED
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
shlq 32, rsi
...
```

Search space - Markov Chain, graph for program space

Figure 3. Search space for the Montgomery multiplication benchmark: 00 and 03 codes are densely connected, whereas expert code is reachable only by an extremely low probability path.



Markov-Chain Monte-Carlo sampling

- Rewrite the cost as a probability distribution

$$p(\mathcal{R}; \mathcal{T}) = \frac{1}{Z} \exp(-\beta \cdot \text{cost}(\mathcal{R}; \mathcal{T}))$$

```
movq rsi, r0
movl ecx, ecx
shrq 32, rsi
andl 0xffffffff, r9d
...
```

```
movq rsi, r0
movl ecx, ecx
UNUSED
andl 0xffffffff, r9d
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
shrq 32, rsi
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
UNUSED
...
```


```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
shlq 32, rsi
...
```

Markov-Chain Monte-Carlo sampling

- Rewrite the cost as a probability distribution

$$p(\mathcal{R}; \mathcal{T}) = \frac{1}{Z} \exp(-\beta \cdot \text{cost}(\mathcal{R}; \mathcal{T}))$$

- Metropolis-Hastings algorithm: Sample from complicated distributions



Picture credit: Colin Carroll

[\[Source\]](#)

```
movq rsi, r0
movl ecx, ecx
shrq 32, rsi
andl 0xffffffff, r9d
...
```

```
movq rsi, r0
movl ecx, ecx
UNUSED
andl 0xffffffff, r9d
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
shrq 32, rsi
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
UNUSED
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
shlq 32, rsi
...
```


Markov-Chain Monte-Carlo sampling

- Rewrite the cost as a probability distribution

$$p(\mathcal{R}; \mathcal{T}) = \frac{1}{Z} \exp(-\beta \cdot \text{cost}(\mathcal{R}; \mathcal{T}))$$


- Metropolis-Hastings algorithm: Sample from complicated distributions
 - Probabilities don't need to be scaled!
 - Based on local decisions in the search space
 - Last rewrite R
 - Proposal rewrite R*

```
movq rsi, r0
movl ecx, ecx
shrq 32, rsi
andl 0xffffffff, r9d
...
```

```
movq rsi, r0
movl ecx, ecx
UNUSED
andl 0xffffffff, r9d
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
shrq 32, rsi
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
UNUSED
...
```



```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
shlq 32, rsi
...
```

Picture credit: Colin Carroll
[\[Source\]](#)

Markov-Chain Monte-Carlo sampling

- Rewrite the cost as a probability distribution

$$p(\mathcal{R}; \mathcal{T}) = \frac{1}{Z} \exp(-\beta \cdot \text{cost}(\mathcal{R}; \mathcal{T}))$$


- Metropolis-Hastings algorithm: Sample from complicated distributions
 - Probabilities don't need to be scaled!
 - Based on local decisions in the search space
 - Last rewrite R
 - Proposal rewrite R*
- Make a proposal from easy distribution q , accept or reject with probability

```
movq rsi, r0
movl ecx, ecx
shrq 32, rsi
andl 0xffffffff, r9d
...
```

```
movq rsi, r0
movl ecx, ecx
UNUSED
andl 0xffffffff, r9d
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
shrq 32, rsi
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
UNUSED
...
```



```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
shlq 32, rsi
...
```

Picture credit: Colin Carroll
[\[Source\]](#)

Markov-Chain Monte-Carlo sampling

- Rewrite the cost as a probability distribution

$$p(\mathcal{R}; \mathcal{T}) = \frac{1}{Z} \exp(-\beta \cdot \text{cost}(\mathcal{R}; \mathcal{T}))$$


- Metropolis-Hastings algorithm: Sample from complicated distributions
 - Probabilities don't need to be scaled!
 - Based on local decisions in the search space
 - Last rewrite R
 - Proposal rewrite R*
- Make a proposal from easy distribution q , accept or reject with probability
 - *If $\text{cost}(\mathcal{R}^*; \mathcal{T}) < \text{cost}(\mathcal{R}; \mathcal{T})$, always accept
 - Else, use the *Metropolis ratio* based on unscaled probabilities

```
movq rsi, r0
movl ecx, ecx
shrq 32, rsi
andl 0xffffffff, r9d
...
```

```
movq rsi, r0
movl ecx, ecx
UNUSED
andl 0xffffffff, r9d
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
shrq 32, rsi
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
UNUSED
...
```



```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
shlq 32, rsi
...
```

Picture credit: Colin Carroll
[\[Source\]](#)

Markov-Chain Monte-Carlo sampling

- In the paper: Synthesis of program from random initialization works well when incremental progress can be made to solution
 - Pretty wild that a program can be synthesized from scratch via sampling!

```
movq rsi, r0
movl ecx, ecx
shrq 32, rsi
andl 0xffffffff, r9d
...
```

```
movq rsi, r0
movl ecx, ecx
UNUSED
andl 0xffffffff, r9d
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
shrq 32, rsi
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
UNUSED
...
```

```
movq rsi, r0
movl ecx, ecx
andl 0xffffffff, r9d
shlq 32, rsi
...
```

Results

Figure 9. Cycling Through 3 Values benchmark.

```
int p21(int x, int a, int b, int c) {  
    return ((-(x == c)) & (a ^ c)) ^  
           ((-(x == a)) & (b ^ c)) ^ c;  
}
```

```
1 # gcc -O3  
2  
3 movl edx, eax  
4 xorl edx, edx  
5 xorl ecx, eax  
6 cmpl esi, edi  
7 sete dl  
8 negl edx  
9 andl edx, eax  
10 xorl edx, edx  
11 xorl ecx, eax  
12 cmpl ecx, edi  
13 sete dl  
14 xorl ecx, esi  
15 negl edx  
16 andl esi, edx  
17 xorl edx, eax
```

```
1 # STOKe  
2  
3 cmpl edi, ecx  
4 cmovel esi, ecx  
5 xorl edi, esi  
6 cmovel edx, ecx  
7 movq rcx, rax
```



Figure 11. Linked List Traversal benchmark.

```
while (head != 0) {  
    head->val *= 2;  
    head = head->next;  
}
```

```
1 # gcc -O3  
2  
3 movq -8(rsp), rdi  
4 .L1:  
5 sall (rdi)  
6 movq 8(rdi), rdi  
7 .L2:  
8 testq rdi, rdi  
9 jne .L1
```

```
1 # STOKe  
2  
3 .L1:  
4 movq -8(rsp), rdi  
5 sall (rdi)  
6 movq 8(rdi), rdi  
7 movq rdi, -8(rsp)  
8 .L2:  
9 movq -8(rsp), rdi  
10 testq rdi, rdi  
11 jne .L1
```



Results

	Speedup ($\times 100\%$)		Runtime (s)	
	gcc/icc -O3	STOKE	Synth.	Opt.
p01	1.60	1.60	0.15	3.05
p02	1.60	1.60	0.16	3.14
p03	1.60	1.60	0.34	3.45
p04	1.60	1.60	2.33	3.55
p05	1.60	1.60	0.47	3.24
p06	1.60	1.60	1.57	6.26
p07	2.00	2.00	1.34	3.10
p08	2.20	2.20	0.63	3.24
p09	1.20	1.20	0.26	3.21
p10	1.80	1.80	7.49	3.61
p11	1.50	1.50	0.87	3.05
p12	1.50	1.50	5.29	3.34
p13	3.25	3.25	0.22	3.08
p14	1.86	1.86	1.43	3.07
p15	2.14	2.14	2.83	3.17
p16	1.80	1.80	6.86	4.62
p17	2.60	2.60	10.65	4.45
p18	2.44	2.50	0.30	4.04
p19	1.93	1.97	-	18.37
p20	1.78	1.78	-	36.72
p21	1.62	1.65	6.97	4.96
p22	3.38	3.41	0.02	4.02
p23	5.53	6.32	0.13	4.36
p24	4.67	4.47	-	48.90
p25	2.17	2.34	3.29	4.43
mont mul	2.84	4.54	319.03	111.64
linked list	1.10	1.09	3.94	8.08
SAXPY	1.82	2.46	10.35	6.66

References & Further Reading

- [1] <http://stoke.stanford.edu/>
- [2] [Stochastic Program Optimization – CACM 2016](#)
- [3] Inference and Metropolis-Hastings in Python: https://colcarroll.github.io/hamiltonian_monte_carlo_talk/bayes_talk.html